

```
!pip install pmdarima


# import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore',category=DeprecationWarning)
import warnings
from statsmodels.tools.sm_exceptions import ConvergenceWarning
warnings.simplefilter('ignore', ConvergenceWarning)
import statsmodels.api as sm
from pylab import rcParams
rcParams['figure.figsize'] = 20,5

from google.colab import drive
drive.mount('/content/drive')

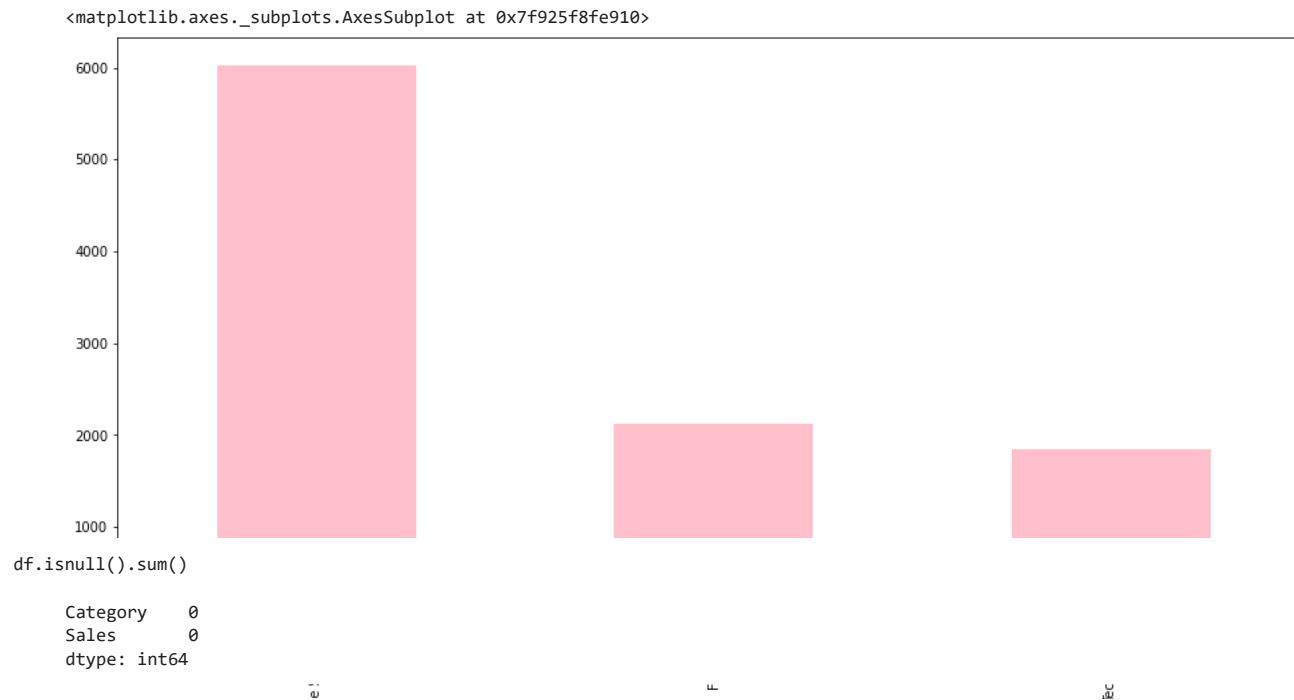
Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
```

## ▼ DATA UNDERSTANDING

```
#READING THE DATASET
df = pd.read_excel("/content/drive/MyDrive/Colab Notebooks/TSF/TimeSeries_Sales_Data_MiniProject.xls",parse_dates=True,index_col='Order Date',usecols=('Category','Sales','Order Date'))
df.head()
```

	Category	Sales	
Order Date			
2016-11-08	Furniture	261.9600	
2016-11-08	Furniture	731.9400	
2016-06-12	Office Supplies	14.6200	
2015-10-11	Furniture	957.5775	
2015-10-11	Office Supplies	22.3680	

```
df['Category'].value_counts().plot(kind='bar',color='pink')
```



## ▼ DATA CATEGORISING FOR ANALYSIS

```
#category wise data frames
one = df.loc[df['Category'] == 'Office Supplies']
two = df.loc[df['Category'] == 'Furniture']
three = df.loc[df['Category'] == 'Technology']
```

```
#selecting only columns required for analysis
offsup = one['Sales']
furn = two['Sales']
tech = three['Sales']
```

## ▼ FUNCTIONS TO USE IN FURTHER MODEL BUILDING AND FORECASTING

```
from statsmodels.tsa.stattools import adfuller
```

```
def checkstationarity(data):
    pvalue = adfuller (data)[1]
    if pvalue < 0.05:
        msg = "pvalue={}. Data is Stationary. Proceed to model building".format(pvalue)
    else:
```

```

    msg = "pvalue={}. Data is not Stationary. Make the data stationary before model building".format(pvalue)
    return msg

from statsmodels.tsa.seasonal import seasonal_decompose

def checkcharacterstics(data):
    dc = seasonal_decompose(data,period=12)
    plt.rcParams.update({'figure.figsize': (16,8)})
    dc.plot().suptitle('Data Decomposition', fontsize=16)

from statsmodels.graphics.tsaplots import plot_acf,plot_pacf

def plotlag(data):
    fig,ax = plt.subplots(1,2,figsize=(20,4))
    plot_acf(data,lags=20,ax=ax[0])
    plot_pacf(data,lags=20,ax=ax[1])
    plt.show()

import statsmodels.api as sm

def ljungboxtest(model):
    pval = sm.stats.acorr_ljungbox(model.resid,lags=[1],return_df=True)['lb_pvalue'].values
    if pval < 0.05:
        print('Reject H0, bad model')
    else:
        print('Accept H0, good model')

def plotpred(data,pred):
    data.plot(label='Actual Dataset')
    pred.plot(label='Predicted Test Set')
    plt.legend()
    plt.show()

def forecasting(model,data):
    global fdata
    forecast = model.predict(start=len(data),end=len(data)+23,dynamic=True)
    fdata = data.append(forecast)
    fdata.plot(label="Current Sales")
    forecast.plot(label="Future Sales")
    plt.legend()
    plt.title("Forecasted Sales for the Upcoming Years")
    plt.show()

```

## ▼ TIME SERIES FORECASTING FOR OFFICE SUPPLIES CATEGORY

```

#CONVERTING INDEX TO DATE TIME FORMAT
offsup.index = pd.to_datetime(offsup.index)


#SORTING THE INDEX
offsup.sort_index(ascending=True,inplace=True)

```

```
#SAMPLING THE DATA MONTHLY FOR DATA REDUCTION
osdata = offsup.resample('MS').sum()
osdata.head()
```

```
Order Date
2014-01-01    4851.080
2014-02-01    1071.724
2014-03-01    8605.879
2014-04-01   11155.074
2014-05-01    7135.624
Freq: MS, Name: Sales, dtype: float64
```

```
#PLOTS TO CHECK TREND,SEASONALITY,RESIDS
checkcharacterstics(osdata)
```



DATA HAS ALL THREE TIME SERIES CHARACTERSTICS TREND, SEASONALITY AND NOISE

```
#CHECKING FOR DATA STATIONARITY
checkstationarity(osdata)
```

```
'pvalue=0.32948727549472817. Data is not Stationary. Make the data stationary before model building'
```


```
osdata_diff = osdata - osdata.shift(1)
osdata_diff.dropna(inplace=True)
checkstationarity(osdata_diff)
```

```
'pvalue=0.00042872624590977896. Data is Stationary. Proceed to model building'
```

1. SINCE DATA IS NON STATIONARY WE CAN MOVE WITH EITHER ARIMA OR SARIMA.
2. SINCE THERE IS SEASONALITY PRESENT I WILL CHOOSE SARIMA MODEL.

```
#CHECKING LAG VALUES
plotlag(osdata_diff)
```

```
/usr/local/lib/python3.8/dist-packages/statsmodels/graphics/tsaplots.py:348: FutureWarning: The default method 'yw' can produce PACF values outside of the [-1,1] interval. After 0.13, the default warnings.warn(
```



```
#FINDING BEST P,Q,D,p,q,d values for building the model
```

```
from pmdarima import auto_arima
m1 = auto_arima(y=osdata, start_p=1, start_q=1, max_p=3, max_q=3, m=12, d=1, D=1, start_P=0, start_Q=0, seasonal=True, error_action='ignore', suppress_warnings=True, trace=True)
print(m1.summary())
```

```
Performing stepwise search to minimize aic
```

```
ARIMA(1,1,1)(0,1,0)[12]      : AIC=730.407, Time=0.07 sec
ARIMA(0,1,0)(0,1,0)[12]      : AIC=740.437, Time=0.03 sec
ARIMA(1,1,0)(1,1,0)[12]      : AIC=726.023, Time=0.12 sec
ARIMA(0,1,1)(0,1,1)[12]      : AIC=717.172, Time=0.28 sec
ARIMA(0,1,1)(0,1,0)[12]      : AIC=728.668, Time=0.05 sec
ARIMA(0,1,1)(1,1,1)[12]      : AIC=718.898, Time=0.45 sec
ARIMA(0,1,1)(0,1,2)[12]      : AIC=718.892, Time=0.87 sec
ARIMA(0,1,1)(1,1,0)[12]      : AIC=718.780, Time=0.26 sec
ARIMA(0,1,1)(1,1,2)[12]      : AIC=725.260, Time=0.54 sec
ARIMA(0,1,0)(0,1,1)[12]      : AIC=730.324, Time=0.08 sec
ARIMA(1,1,1)(0,1,1)[12]      : AIC=722.788, Time=0.17 sec
ARIMA(0,1,2)(0,1,1)[12]      : AIC=718.277, Time=0.62 sec
ARIMA(1,1,0)(0,1,1)[12]      : AIC=722.473, Time=0.37 sec
ARIMA(1,1,2)(0,1,1)[12]      : AIC=723.414, Time=0.98 sec
ARIMA(0,1,1)(0,1,1)[12] intercept : AIC=723.739, Time=0.15 sec
```

```
Best model: ARIMA(0,1,1)(0,1,1)[12]
```

```
Total fit time: 5.070 seconds
```

```
SARIMAX Results
```

```

=====
Dep. Variable:          y      No. Observations:          48
Model:                SARIMAX(0, 1, 1)x(0, 1, 1, 12)  Log Likelihood          -355.586
Date:                  Tue, 31 Jan 2023              AIC              717.172
Time:                  14:16:16                      BIC              721.838
Sample:                01-01-2014                   HQIC              718.783
                   - 12-01-2017
Covariance Type:      opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ma.L1          -0.6782      0.117      -5.811      0.000      -0.907      -0.449
ma.S.L12        -0.6428      0.285      -2.256      0.024      -1.201      -0.084
sigma2          3.679e+07  4.94e-09  7.45e+15      0.000      3.68e+07  3.68e+07
=====
Ljung-Box (L1) (Q):                0.01  Jarque-Bera (JB):                23.75
Prob(Q):                           0.93  Prob(JB):                      0.00
Heteroskedasticity (H):             2.76  Skew:                          1.30
Prob(H) (two-sided):               0.09  Kurtosis:                      6.08
=====

```

## Warnings:

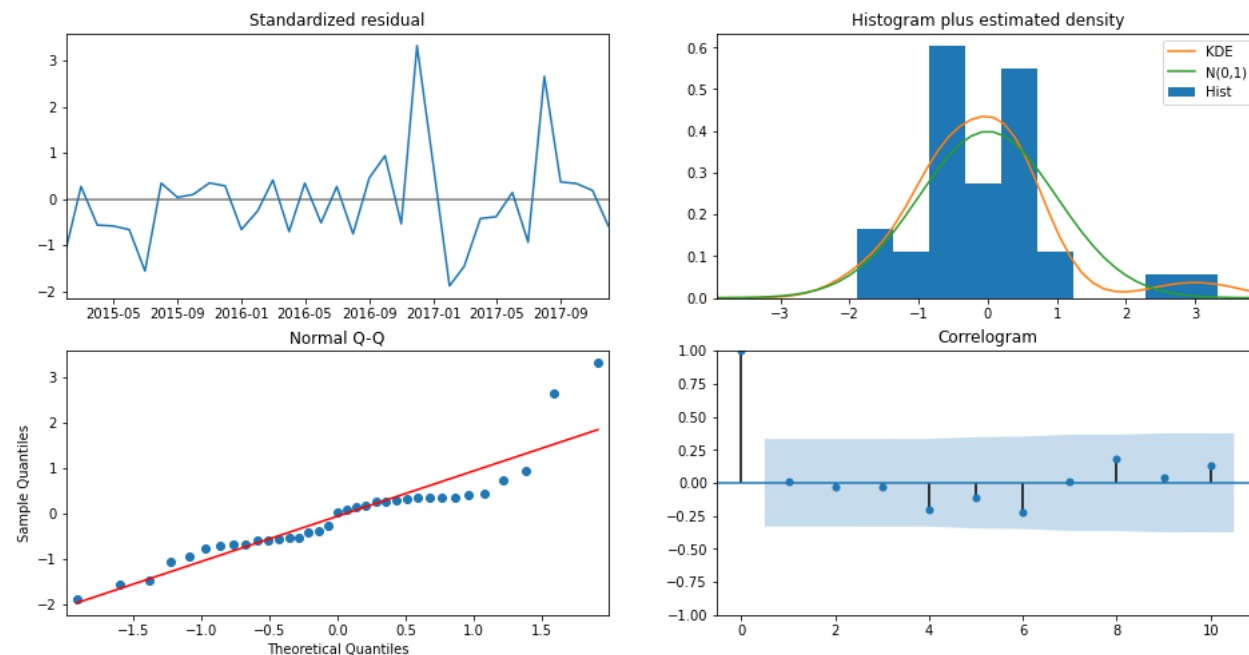
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 2.64e+31. Standard errors may be unstable.

#CHECKING FOR MODEL DIAGNOSTICS BEFORE BULIDING

m1.plot\_diagnostics()

plt.show()



1. Standardised Residuals : There are obvious patterns in the residuals with no uniform mean and variance.
2. KDE Curve : Does not show normal distributions, data is skewed.
3. Normal Q-Q : Most of the data points are not on the straight line/normal reference line.

```
split1 = int(0.8 * len(osdata))
train1 = osdata[:split1]
test1 = osdata[split1:]
```

#BUILDING THE SARIMA MODEL FOR OFFICE SUPPLIES

```
os_model = sm.tsa.statespace.SARIMAX(train1, order=(1,1,1), seasonal_order=(0,1,1,12)).fit()
print(os_model.summary())
```

```

SARIMAX Results
=====
Dep. Variable:          Sales      No. Observations:          38
Model:              SARIMAX(1, 1, 1)x(0, 1, 1, 12)      Log Likelihood          -253.947
Date:                Tue, 31 Jan 2023      AIC              515.894
Time:                14:57:14              BIC              520.770
Sample:              01-01-2014      HQIC             517.247
                  - 02-01-2017
Covariance Type:      opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1           0.0075      0.415      0.018      0.986      -0.805      0.820
ma.L1          -0.6175      0.320     -1.930      0.054      -1.245      0.010
ma.S.L12       -0.4241      0.741     -0.573      0.567      -1.876      1.027
sigma2         5.149e+07   8.66e-09   5.94e+15      0.000      5.15e+07   5.15e+07
=====
Ljung-Box (L1) (Q):           0.45      Jarque-Bera (JB):          15.68
Prob(Q):                   0.50      Prob(JB):              0.00
Heteroskedasticity (H):       3.60      Skew:                  1.10
Prob(H) (two-sided):         0.09      Kurtosis:              6.19
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 4.75e+31. Standard errors may be unstable.

/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/statespace/sarimax.py:866: UserWarning: Too few observations to estimate starting parameters for seasonal ARMA. All parameters except for va warn('Too few observations to estimate starting parameters%s.'


#LJUNG BOX TEST

```
ljungoxtest(os_model)
```

Accept  $H_0$ , good model

#PERFORMANCE METRICS and PREDICTED DATA

```
from sklearn.metrics import mean_squared_error, mean_absolute_percentage_error
pred1 = os_model.predict(start=len(train1), end=len(osdata)-1, dynamic=True)
mse1 = mean_squared_error(test1, pred1)
mape1 = mean_absolute_percentage_error(test1, pred1)
plotpred(osdata, pred1)
```



```
#FORECASTING
forecasting(os_model,osdata)
```



```
/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/statespace/kalman_filter.py:2290: ValueWarning: Dynamic prediction specified to begin during out-of-sample forecasting period, and so has no
warn('Dynamic prediction specified to begin during'
```

#### INTERPRETATIONS:

1. SARIMA MODEL WAS USED HERE FOR FORECASTING DUE TO THE PRESENCE OF SEASONALITY.
2. DATA WAS NOT STATIONARY IN THIS CATEGORY.
3. PATTERNS ARE QUIET SIMILAR TO THE PREVIOUS YEAR SALES RECORD.
4. UPWARD TREND CAN BE SEEN IN THE OVERALL DATA AND THE DATA FORECASTED.
5. YEAR END SALES IN OFFICE SUPPLY CATEGORY ARE COMPARATIVELY MORE.

### TIME SERIES FORECASTING FOR FURNITURE CATEGORY

```
#CONVERTING INDEX TO DATE TIME FORMAT
furn.index = pd.to_datetime(furn.index)
```


```
#SORTING THE INDEX
furn.sort_index(ascending=True,inplace=True)
```

```
#SAMPLING THE DATA MONTHLY FOR DATA REDUCTION
furndata = furn.resample('MS').sum()
furndata.head()
```

```
Order Date
2014-01-01    6242.525
2014-02-01    1839.658
2014-03-01   14573.956
2014-04-01    7944.837
2014-05-01    6912.787
Freq: MS, Name: Sales, dtype: float64
```

```
#SMOOTHING
from statsmodels.tsa.api import ExponentialSmoothing
df2_smo = ExponentialSmoothing(furndata,trend='add',seasonal='add',seasonal_periods=12).fit()
furndata = df2_smo.fittedvalues
```

```
#PLOTS TO CHECK TREND,SEASONALITY,RESIDS
checkcharacterstics(furndata)
```



DATA HAS ALL THREE TIME SERIES CHARACTERISTICS TREND, SEASONALITY AND NOISE



```
#CHECKING FOR DATA STATIONARITY
checkstationarity(furndata)
```

```
'pvalue=9.513079933718842e-05. Data is Stationary. Proceed to model building'
```

1. SINCE DATA IS STATIONARY WE CAN MOVE WITH ANY OF THE EITHER ARMA/ARIMA/SARIMA.
2. SINCE THERE IS SEASONALITY PRESENT I WILL CHOOSE SARIMA MODEL.

```
#CHECKING LAG VALUES
plotlag(furndata)
```

```
/usr/local/lib/python3.8/dist-packages/statsmodels/graphics/tsaplots.py:348: FutureWarning: The default method 'yw' can produce PACF values outside of the [-1,1] interval. After 0.13, the default warnings.warn()
```



```
#FINDING BEST P,Q,D,p,q,d values for building the model
```

```
m2 = auto_arma(y=furndata,start_p=1,start_q=1,max_p=9,max_q=2,m=12,d=0,D=0,start_P=0,start_Q=0,seasonal=True,error_action='ignore', suppress_warnings=True,trace=True)
```

```
print(m2.summary())
```

```
Performing stepwise search to minimize aic
```

```
ARIMA(1,0,1)(0,0,0)[12] intercept : AIC=1008.144, Time=0.04 sec
ARIMA(0,0,0)(0,0,0)[12] intercept : AIC=1008.133, Time=0.02 sec
ARIMA(1,0,0)(1,0,0)[12] intercept : AIC=975.883, Time=0.31 sec
ARIMA(0,0,1)(0,0,1)[12] intercept : AIC=989.458, Time=0.08 sec
ARIMA(0,0,0)(0,0,0)[12] intercept : AIC=1074.445, Time=0.01 sec
ARIMA(1,0,0)(0,0,0)[12] intercept : AIC=1006.161, Time=0.03 sec
ARIMA(1,0,0)(2,0,0)[12] intercept : AIC=990.474, Time=0.22 sec
ARIMA(1,0,0)(1,0,1)[12] intercept : AIC=985.751, Time=0.14 sec
ARIMA(1,0,0)(0,0,1)[12] intercept : AIC=987.766, Time=0.14 sec
ARIMA(1,0,0)(2,0,1)[12] intercept : AIC=inf, Time=2.67 sec
ARIMA(0,0,0)(1,0,0)[12] intercept : AIC=1003.891, Time=0.13 sec
ARIMA(2,0,0)(1,0,0)[12] intercept : AIC=989.294, Time=0.27 sec
ARIMA(1,0,1)(1,0,0)[12] intercept : AIC=inf, Time=1.38 sec
ARIMA(0,0,1)(1,0,0)[12] intercept : AIC=1001.164, Time=0.14 sec
ARIMA(2,0,1)(1,0,0)[12] intercept : AIC=inf, Time=0.91 sec
ARIMA(1,0,0)(1,0,0)[12] intercept : AIC=inf, Time=0.45 sec
```

```
Best model: ARIMA(1,0,0)(1,0,0)[12] intercept
```

```
Total fit time: 6.980 seconds
```

#### SARIMAX Results

```
=====
Dep. Variable: y No. Observations: 48
Model: SARIMAX(1, 0, 0)x(1, 0, 0, 12) Log Likelihood -483.941
Date: Tue, 31 Jan 2023 AIC 975.883
Time: 15:03:21 BIC 983.367
Sample: 01-01-2014 HQIC 978.711
- 12-01-2017
```

```
Covariance Type: opg
```

```
=====
coef std err z P>|z| [0.025 0.975]
-----
intercept 9373.4714 1560.511 6.007 0.000 6314.927 1.24e+04
ar.L1 -0.2183 0.192 -1.135 0.256 -0.595 0.159
ar.S.L12 0.5494 0.105 5.244 0.000 0.344 0.755
sigma2 2.604e+07 0.227 1.15e+08 0.000 2.6e+07 2.6e+07
=====
```

```
Ljung-Box (L1) (Q): 12.36 Jarque-Bera (JB): 1.38
Prob(Q): 0.00 Prob(JB): 0.50
Heteroskedasticity (H): 0.44 Skew: 0.40
Prob(H) (two-sided): 0.11 Kurtosis: 2.78
=====
```

#### Warnings:

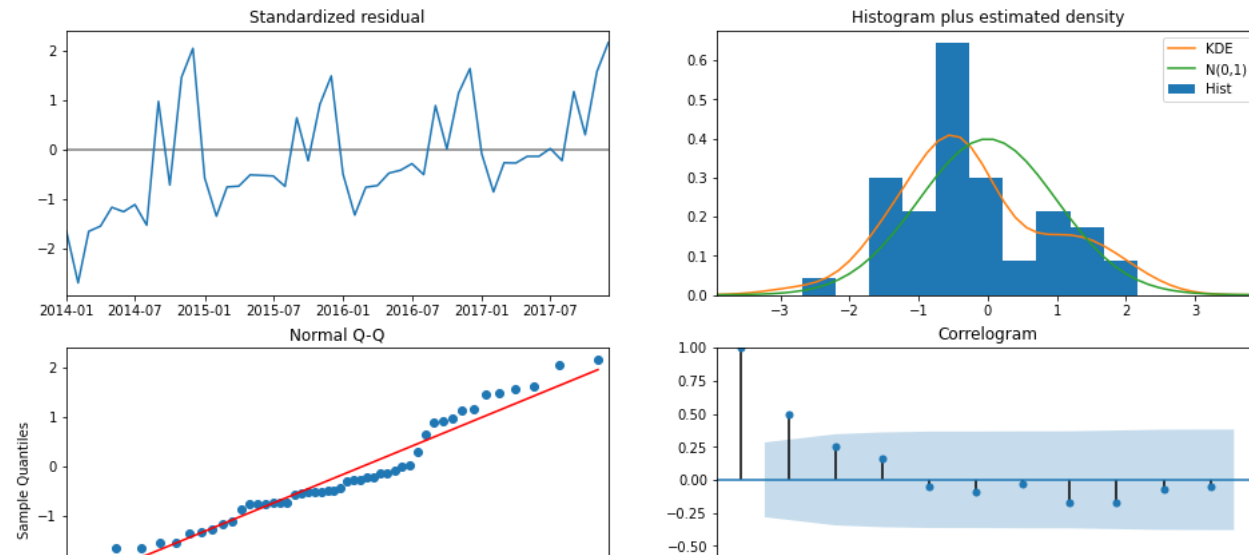
```
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

```
[2] Covariance matrix is singular or near-singular, with condition number 4.8e+23. Standard errors may be unstable.
```

```
#CHECKING FOR MODEL DIAGNOSTICS BEFORE BULIDING
```

```
m2.plot_diagnostics()
```

```
plt.show()
```



1. Standardised Residuals : There are no obvious patterns in the residuals.
2. KDE Curve : Shows normal distributions, data is skewed.
3. Normal Q-Q : Most of the data points are on the straight line/normal reference line.

```
split2 = int(0.8 * len(furndata))
train2 = furndata[:split2]
test2 = furndata[split2:]
```

#BUILDING THE SARIMA MODEL FOR OFFICE SUPPLIES

```
furn_model = sm.tsa.statespace.SARIMAX(train2, order=(1,0,1), seasonal_order=(1,0,0,12)).fit()
print(furn_model.summary())
```

```
/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/statespace/sarimax.py:997: UserWarning: Non-stationary starting seasonal autoregressive Using zeros as starting parameters.
warn('Non-stationary starting seasonal autoregressive')
```

SARIMAX Results

```
=====
Dep. Variable:          y      No. Observations:          38
Model:          SARIMAX(1, 0, 1)x(1, 0, [], 12)      Log Likelihood          -358.764
Date:              Tue, 31 Jan 2023      AIC              725.529
Time:              15:03:26      BIC              732.079
Sample:              01-01-2014      HQIC             727.859
              - 02-01-2017
```

Covariance Type: opg

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.9998      0.133       7.500      0.000      0.739      1.261
ma.L1         -0.9994      0.259      -3.854      0.000     -1.508     -0.491
ar.S.L12        0.9878      0.010     102.698      0.000      0.969      1.007
sigma2       5.302e+06     1.87e-08     2.83e+14      0.000     5.3e+06     5.3e+06
=====
```

```
Ljung-Box (L1) (Q):              4.88      Jarque-Bera (JB):              14.15
```

```

Prob(Q):                0.03   Prob(JB):                0.00
Heteroskedasticity (H):  0.53   Skew:                  1.33
Prob(H) (two-sided):    0.27   Kurtosis:          4.38
=====

```

## Warnings:

```

[1] Covariance matrix calculated using the outer product of gradients (complex-step).
[2] Covariance matrix is singular or near-singular, with condition number 6.51e+30. Standard errors may be unstable.

```

## #LJUNG BOX TEST

```
ljungtest(furn_model)
```


```
Reject H0, bad model
```

## #PERFORMANCE METRICS and PREDICTED DATA

```

pred2 = os_model.predict(start=len(train2),end=len(furndata)-1,dynamic=True)
mse2 = mean_squared_error(test2,pred2)
mape2 = mean_absolute_percentage_error(test2,pred2)
plotpred(furndata,pred2)


```



## #FORECASTING

```
forecasting(furn_model,furndata)
```

```
/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/statespace/kalman_filter.py:2290: ValueWarning: Dynamic prediction specified to begin during out-of-sample forecasting period, and so has no warn('Dynamic prediction specified to begin during'
```



#### INTERPRETATIONS:

1. SARIMA MODEL WAS USED HERE FOR FORECASTING DUE TO THE PRESENCE OF SEASONALITY.
2. DATA WAS STATIONARY IN THIS CATEGORY.
3. PATTERNS ARE VERY SIMILAR TO THE PREVIOUS YEAR SALES RECORD.
4. VERY SLIGHT UPWARD TREND CAN BE SEEN IN THE OVERALL DATA AND THE DATA FORECASTED.
5. YEAR END SALES IN FURNITURE SALES ARE COMPARATIVELY MORE.
6. DIPS IN SALES CAN BE SEEN 1ST QUARTER.

## ▼ TIME SERIES FORECASTING FOR TECHNOLOGY CATEGORY

```
#CONVERTING INDEX TO DATE TIME FORMAT
tech.index = pd.to_datetime(tech.index)
```

```
#SORTING THE INDEX
tech.sort_index(ascending=True,inplace=True)
```


```
#SAMPLING THE DATA MONTHLY FOR DATA REDUCTION
techdata = tech.resample('MS').sum()
techdata.head()
```

```
Order Date
2014-01-01    3143.290
2014-02-01    1608.510
2014-03-01   32511.174
2014-04-01    9195.434
```

```
2014-05-01    9599.876
Freq: MS, Name: Sales, dtype: float64
```

```
#SMOOTHING
from statsmodels.tsa.api import ExponentialSmoothing
df_smo = ExponentialSmoothing(techdata,trend='add',seasonal='add',seasonal_periods=12).fit()
techdata = df_smo.fittedvalues
```

```
#PLOTS TO CHECK TREND,SEASONALITY,RESIDS
checkcharacterstics(techdata)
```



DATA HAS ALL THREE TIME SERIES CHARACTERSTICS TREND, SEASONALITY AND NOISE

```
#CHECKING FOR DATA STATIONARITY
checkstationarity(techdata)
```

```
'pvalue=0.9852373582089293. Data is not Stationary. Make the data stationary before model building'
```

```
techdata_diff = techdata - techdata.shift(1)
techdata_diff.dropna(inplace=True)
checkstationarity(techdata_diff)
```


```
'pvalue=1.5818270748199294e-16. Data is Stationary. Proceed to model building'
```

1. SINCE DATA IS NOT STATIONARY WE CAN MOVE WITH ANY OF THE EITHER ARMA/ARIMA/SARIMA.
2. SINCE THERE IS SEASONALITY PRESENT I WILL CHOOSE SARIMA MODEL.

```
#CHECKING LAG VALUES
```

```
plotlag(techdata_diff)
```

```
/usr/local/lib/python3.8/dist-packages/statsmodels/graphics/tsaplots.py:348: FutureWarning: The default method 'yw' can produce PACF values outside of the [-1,1] interval. After 0.13, the default warnings.warn()
```



```
#FINDING BEST P,Q,D,p,q,d values for building the model
```

```
m3 = auto_arma(y=techdata,start_p=1,start_q=1,max_p=14,max_q=3,m=12,d=1,D=1,start_P=0,start_Q=0,seasonal=True,error_action='ignore',suppress_warnings=True,trace=True)
print(m3.summary())
```

```
Performing stepwise search to minimize aic
```

```
ARIMA(1,1,1)(0,1,0)[12]      : AIC=602.198, Time=0.07 sec
ARIMA(0,1,0)(0,1,0)[12]      : AIC=598.230, Time=0.03 sec
ARIMA(1,1,0)(1,1,0)[12]      : AIC=601.606, Time=0.09 sec
ARIMA(0,1,1)(0,1,1)[12]      : AIC=601.604, Time=0.08 sec
ARIMA(0,1,0)(1,1,0)[12]      : AIC=599.802, Time=0.05 sec
ARIMA(0,1,0)(0,1,1)[12]      : AIC=599.680, Time=0.05 sec
ARIMA(0,1,0)(1,1,1)[12]      : AIC=inf, Time=0.38 sec
ARIMA(1,1,0)(0,1,0)[12]      : AIC=600.176, Time=0.07 sec
ARIMA(0,1,1)(0,1,0)[12]      : AIC=600.193, Time=0.04 sec
ARIMA(0,1,0)(0,1,0)[12] intercept : AIC=599.435, Time=0.02 sec
```

```
Best model: ARIMA(0,1,0)(0,1,0)[12]
```

```
Total fit time: 0.899 seconds
```

```
SARIMAX Results
```

```
=====
Dep. Variable:          y      No. Observations:          48
Model:          SARIMAX(0, 1, 0)x(0, 1, 0, 12)      Log Likelihood          -298.115
Date:          Tue, 31 Jan 2023      AIC          598.230
Time:          14:16:30      BIC          599.785
Sample:          01-01-2014      HQIC          598.767
              - 12-01-2017
```

```
Covariance Type:          opg
```

```
=====
coef      std err          z      P>|z|      [0.025      0.975]
-----

```



```

sigma2      1.462e+06   3.66e+05   3.993   0.000   7.44e+05   2.18e+06
=====
Ljung-Box (L1) (Q):           0.00   Jarque-Bera (JB):           1.09
Prob(Q):                     0.98   Prob(JB):                 0.58
Heteroskedasticity (H):       1.45   Skew:                    -0.43
Prob(H) (two-sided):          0.53   Kurtosis:                 3.07
=====

```

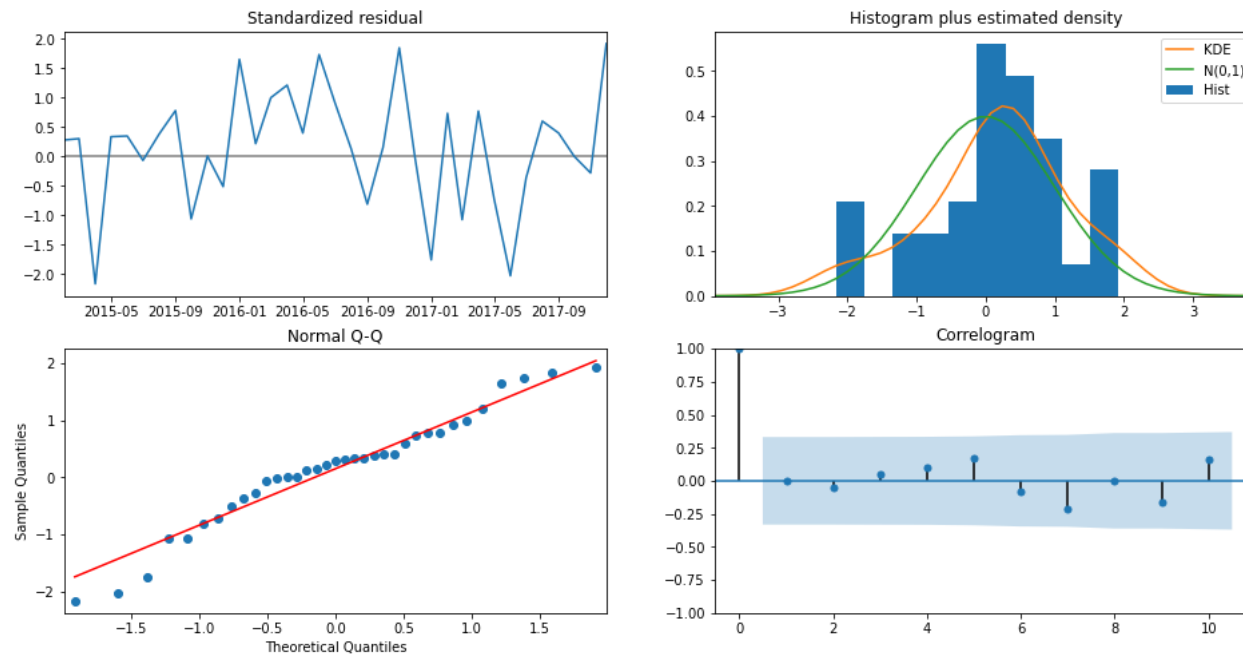
Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

#CHECKING FOR MODEL DIAGNOSTICS BEFORE BULIDING

m3.plot\_diagnostics()

plt.show()



1. Standardised Residuals : Thereare obvious patterns in the residuals with no uniform mean and variance.
2. KDE Curve : Dosent show normal distributions, data is skewed.
3. Normal Q-Q : Most of the data points are not on the straight line/normal reference line.

```

split3 = int(0.8 * len(techdata))
train3 = techdata[:split3]
test3 = techdata[split3:]

```

```
#BUILDING THE SARIMA MODEL FOR OFFICE SUPPLIES
```

```
tech_model = sm.tsa.statespace.SARIMAX(train3,order=(1,1,1),seasonal_order=(0,1,0,12)).fit()
```

```
print(tech_model.summary())
```

```

SARIMAX Results
=====
Dep. Variable:          y      No. Observations:          38
Model:          SARIMAX(1, 1, 1)x(0, 1, [], 12)  Log Likelihood          -212.577
Date:              Tue, 31 Jan 2023              AIC              431.154
Time:              15:06:07                      BIC              434.811
Sample:           01-01-2014                    HQIC              432.169
              - 02-01-2017
Covariance Type:          opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.0709        7.925        0.009        0.993       -15.463       15.604
ma.L1         -0.0871        7.867       -0.011        0.991       -15.507       15.333
sigma2        1.283e+06    3.56e+05        3.603        0.000     5.85e+05    1.98e+06
=====
Ljung-Box (L1) (Q):          0.00  Jarque-Bera (JB):          1.66
Prob(Q):          0.98  Prob(JB):          0.44
Heteroskedasticity (H):        1.50  Skew:          -0.60
Prob(H) (two-sided):        0.58  Kurtosis:          3.40
=====

```

```
Warnings:
```

```
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

```
/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/statespace/sarimax.py:966: UserWarning: Non-stationary starting autoregressive parameters found. Using zeros as starting parameters.
```

```
warn('Non-stationary starting autoregressive parameters')
```

```
/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/statespace/sarimax.py:978: UserWarning: Non-invertible starting MA parameters found. Using zeros as starting parameters.
```

```
warn('Non-invertible starting MA parameters found.')
```

```
#LJUNG BOX TEST
```

```
ljungboxtest(tech_model)
```

```
Accept H0, good model
```


```
#PERFORMACE METRICS and PREDICTED DATA
```

```
pred3 = tech_model.predict(start=len(train3),end=len(techdata)-1,dynamic=True)
```

```
mse3 = mean_squared_error(test3,pred3)
```

```
mape3 = mean_absolute_percentage_error(test3,pred3)
```


```
plotpred(techdata,pred3)
```



```
#FORECASTING
forecasting(tech_model,techdata)
```

```
/usr/local/lib/python3.8/dist-packages/statsmodels/tsa/statespace/kalman_filter.py:2290: ValueWarning: Dynamic prediction specified to begin during out-of-sample forecasting period, and so has no warn('Dynamic prediction specified to begin during')
```

Forecasted Sales for the Upcoming Years



#### INTERPRETATIONS:

1. SARIMA MODEL WAS USED HERE FOR FORECASTING DUE TO THE PRESENCE OF SEASONALITY.
2. DATA WAS NOT STATIONARY IN THIS CATEGORY.
3. PATTERNS ARE QUIET SIMILAR TO THE PREVIOUS YEAR SALES RECORD.
4. STRONG UPWARD TREND CAN BE SEEN IN THE OVERALL DATA AND THE DATA FORECASTED.
5. YEAR END SALES IN TECH SUPPLY CATEGORY ARE COMPARITIVELY MORE.

```
frame = {'CATEGORY': ['OFFICE SUPPLIES', 'FURNITURE', 'TECHNOLOGY'], 'RMSE': [np.sqrt(mse1), np.sqrt(mse2), np.sqrt(mse3)], 'MAPE': [mape1, mape2, mape3]}
table = pd.DataFrame(frame)
table
```

	CATEGORY	RMSE	MAPE
0	OFFICE SUPPLIES	6586.551225	0.327788
1	FURNITURE	6140.841973	0.383515
2	TECHNOLOGY	2721.212192	0.113917

MAPE is defined as the percentage of the average of absolute difference between forecasted values and true values, divided by true value.

The lower the MAPE, the better the model is.

## INNOVATION


```
techdata = pd.DataFrame(fdata)
techdata
```

	0
2014-01-01	6729.032286
2014-02-01	4051.774224
2014-03-01	18015.409966
2014-04-01	10361.914810
2014-05-01	14801.616282
...	...
2019-08-01	34917.944206
2019-09-01	47352.872265
2019-10-01	41075.208297
2019-11-01	51650.875246
2019-12-01	50634.243836

72 rows × 1 columns

```
techdata[1] = list(techdata.index)
```

```
techdata
```

	Technology_Sales	Date	
2014-01-01	6729.032286	2014-01-01	
2014-02-01	4051.774224	2014-02-01	
2014-03-01	18015.409966	2014-03-01	
2014-04-01	10361.914810	2014-04-01	
2014-05-01	14801.616282	2014-05-01	
...	...	...	
2019-08-01	34917.944206	2019-08-01	
2019-09-01	47352.872265	2019-09-01	
2019-10-01	41075.208297	2019-10-01	
2019-11-01	51650.875246	2019-11-01	

```

techdata.rename(columns={0: 'Technology_Sales', 1: 'Date'}, inplace=True)
72 rows x 2 columns
techdata.to_csv('Technology Sales.csv')

from matplotlib.animation import FuncAnimation
from itertools import count
x=[]
y=[]
fig,ax = plt.subplots()
ax.plot(x,y)
plt.style.use("ggplot")
counter=count(0,1)

def update(i):
    idx=next(counter)
    x.append(techdata.loc[idx,techdata['Date']])
    y.append(techdata.loc[idx,techdata['Technology_Sales']])
    plt.cla()
    ax.plot(x,y)

ani=FuncAnimation(fig=fig,func=update,interval=200,frames = len(x) + 1)
plt.show()

```


Traceback (most recent call last):

```
File "/usr/local/lib/python3.8/dist-packages/matplotlib/cbook/__init__.py", line 196, in process
    func(*args, **kwargs)
File "/usr/local/lib/python3.8/dist-packages/matplotlib/animation.py", line 951, in _start
    self._init_draw()
File "/usr/local/lib/python3.8/dist-packages/matplotlib/animation.py", line 1743, in _init_draw
    self._draw_frame(next(self.new_frame_seq()))
File "/usr/local/lib/python3.8/dist-packages/matplotlib/animation.py", line 1766, in _draw_frame
    self._drawn_artists = self._func(framedata, *self._args)
File "<ipython-input-410-b93ab201a488>", line 12, in update
    x.append(techdata.loc[idx,techdata['Date']])
File "/usr/local/lib/python3.8/dist-packages/pandas/core/indexing.py", line 925, in __getitem__
    return self._getitem_tuple(key)
File "/usr/local/lib/python3.8/dist-packages/pandas/core/indexing.py", line 1100, in _getitem_tuple
    return self._getitem_lowerdim(tup)
File "/usr/local/lib/python3.8/dist-packages/pandas/core/indexing.py", line 838, in _getitem_lowerdim
    section = self._getitem_axis(key, axis=i)
File "/usr/local/lib/python3.8/dist-packages/pandas/core/indexing.py", line 1164, in _getitem_axis
    return self._get_label(key, axis=axis)
File "/usr/local/lib/python3.8/dist-packages/pandas/core/indexing.py", line 1113, in _get_label
    return self.obj.xs(label, axis=axis)
File "/usr/local/lib/python3.8/dist-packages/pandas/core/generic.py", line 3776, in xs
    loc = index.get_loc(key)
File "/usr/local/lib/python3.8/dist-packages/pandas/core/indexes/datetimes.py", line 700, in get_loc
    raise KeyError(key)
KeyError: 0
```



```
f1 = pd.read_csv("/content/Office Suuply Sales.csv",parse_dates=True,usecols=('Office_Supplies_Sales','Date'),index_col='Date')
f2 = pd.read_csv("/content/Furniture Sales.csv",parse_dates=True,usecols=('Furniture_Sales','Date'),index_col='Date')
f3 = pd.read_csv("/content/Technology Sales.csv",parse_dates=True,usecols=('Technology_Sales','Date'),index_col='Date')
```

f3

Technology\_Sales 

Date

2014-01-01	6729.032286
2014-02-01	4051.774224
2014-03-01	18015.409966
2014-04-01	10361.914810
2014-05-01	14801.616282

f1['Furniture\_Sales'] = f2.Furniture\_Sales


f1['Technology\_Sales'] = f3.Technology\_Sales

2019-09-01	47352.872265
------------	--------------

category\_forecasts = f1.copy()

...	...
-----	-----

category\_forecasts

Office\_Supplies\_Sales Furniture\_Sales Technology\_Sales 

Date

2014-01-01	4851.080000	6790.384086	6729.032286
2014-02-01	1071.724000	2880.555398	4051.774224
2014-03-01	8605.879000	10091.659067	18015.409966
2014-04-01	11155.074000	9156.845289	10361.914810
2014-05-01	7135.624000	11687.553932	14801.616282
...	...	...	...
2019-08-01	29869.816649	10376.759128	34917.944206
2019-09-01	41815.848384	25988.638943	47352.872265
2019-10-01	32883.562624	12937.631855	41075.208297
2019-11-01	40620.064180	28212.634448	51650.875246
2019-12-01	48467.427362	30292.353643	50634.243836


72 rows × 3 columns

category\_forecasts.to\_csv('category\_forecasts.csv')


category\_forecasts.plot()

plt.axvline(x='2018-01-01',c='red')

&lt;matplotlib.lines.Line2D at 0x7f925f8fbd60&gt;



```
import plotly.express as px
fig = px.line(category_forecasts, x=category_forecasts.index, y=category_forecasts.columns)
fig.show()
```





---

✓ 0s completed at 11:14 PM

